



Producer Scripts – API to Pub/Sub Component

GCP-Ready Producer Scripts for Secure API Fetch with GSM Key & Pub/Sub Publishing

Overview

This component provides Python scripts to securely fetch data from the hospital API endpoints (protected by API keys in GSM) and publish to Google Cloud Pub/Sub topics. It is designed for robust, auditable, and cloud-native data streaming pipelines.

- **Language:** Python 3.11+
- **Security:** Google Secret Manager (API Key)
- **Transport:** HTTPS (requests)
- **Target:** Google Cloud Pub/Sub (Streaming)
- **Usage:** Fits into any modern GCP streaming or batch ETL flow.



VM Setup, Firewall & Permissions (GCP & Local)

1. VM Creation (GCP)

- Use a **Debian/Ubuntu** image with at least 2 vCPUs and 4GB RAM for production streaming.
- Assign a **service account** with the following roles:
- `roles/secretmanager.secretAccessor`
- `roles/pubsub.publisher`
- Enable required APIs:

```
gcloud services enable secretmanager.googleapis.com
pubsub.googleapis.com
```

2. Firewall Rules

- **Outbound (egress):** Open by default. Needed for:
- HTTPS (443) to Google APIs (GSM, Pub/Sub)
- HTTP (80) if your APIs run on HTTP (dev/test)
- **Inbound (ingress):**
- **SSH (22):** Open only to your trusted IPs for admin
- **No public inbound ports needed** for Producer unless you expose an API/UI (not typical).

3. Local VM/Server (Non-GCP)

- Ensure outbound access to:
- `*.googleapis.com:443` (GCP APIs)
- API VM IP/port (e.g., `:5001`, `:5002`, `:5003`)
- If using a proxy/firewall, whitelist the above.

4. Permissions Checklist

- **Producer Service Account** must have:
- `roles/secretmanager.secretAccessor`
- `roles/pubsub.publisher` for your target Pub/Sub topics
- **API VM Firewall:** Allow traffic **from Producer VM's IP** to ports 5001–5003.

Summary Table

Resource/Port	Direction	Rule/Permission	Why
SSH (22)	Inbound	Open to admin IPs	Remote VM admin
GSM, Pub/Sub (443)	Outbound	Open (default in GCP)	Cloud API access (GSM, Pub/Sub)
Hospital API (5001-3)	Outbound	Open to API VM IPs	Fetch data from API
SecretManager Access	-	<code>roles/secretmanager.secretAccessor</code>	Fetch API keys
Pub/Sub Publisher	-	<code>roles/pubsub.publisher</code>	Publish messages

```
producer/
├─ producer_patients.py
├─ producer_appointments.py
├─ producer_lab_results.py
├─ shared/
│   ├─ config_util.py          # Loads config.yaml
│   ├─ fetch_api_key.py        # Fetches key from GSM
│   └─ config.yaml             # Stores API URLs and Pub/Sub topic names
└─ requirements.txt
```

Setup & Quickstart

Security Best Practice: Always export `API_KEY_SECRET_ID` (the full GSM secret resource path) as an environment variable on your Producer VM **before running any script**. This ensures your code fetches the API key securely from Google Secret Manager and never from local files or code.

1. Clone & Prepare Python Environment

```
cd producer
python3.11 -m venv venv
```

```
source venv/bin/activate
pip install -r requirements.txt
```

2. Configure API URLs, Topics, and GSM Secret

- Edit `shared/config.yaml`:

```
api:
  patients_url: "http://<API_VM_IP>:5001/patients"
  appointments_url: "http://<API_VM_IP>:5002/appointments"
  lab_results_url: "http://<API_VM_IP>:5003/lab_results"
pubsub:
  patients_topic: "projects/<PROJECT_ID>/topics/patients-stream"
  appointments_topic: "projects/<PROJECT_ID>/topics/appointments-stream"
  lab_results_topic: "projects/<PROJECT_ID>/topics/lab-results-stream"
```

- Set your **API Key GSM secret path** as an environment variable:

```
export API_KEY_SECRET_ID="projects/<PROJECT_ID>/secrets/hospital-api-
key/versions/latest"
```

3. Run Producer Scripts

```
python producer_patients.py
python producer_appointments.py
python producer_lab_results.py
```

How Secure API Key Fetch Works

- API key is never hardcoded—always read securely from GSM using the env variable.
- Code in `shared/fetch_api_key.py`:

```
import os
from google.cloud import secretmanager

def get_api_key():
    secret_id = os.environ["API_KEY_SECRET_ID"]
    client = secretmanager.SecretManagerServiceClient()
    response = client.access_secret_version(request={"name": secret_id})
    return response.payload.data.decode('UTF-8')
```

- All requests to API endpoints use the `"x-api-key"` header.

End-to-End Data Flow

```
graph TD
  API_VM["Hospital API VM (Gunicorn/Flask)"]
  PRODUCER["Producer VM (This Component)"]
  PUBSUB["Google Pub/Sub"]
  API_VM-->|Secure API Key|PRODUCER
  PRODUCER-->|publish JSON|PUBSUB
```

requirements.txt

This covers **all packages imported in every script in this component**:

```
requests
google-cloud-secret-manager
google-cloud-pubsub
pyyaml
```

Install with: `pip install -r requirements.txt`

Example Usage (Python)

```
# From producer_patients.py
from shared.config_util import load_config
from shared.fetch_api_key import get_api_key
import requests

config = load_config()
API_URL = config['api']['patients_url']
API_KEY = get_api_key()
resp = requests.get(API_URL, headers={"x-api-key": API_KEY})
data = resp.json()["patients"]
```

Security & IAM Checklist

- Producer VM/service account must have:
 - `roles/secretmanager.secretAccessor`
 - `roles/pubsub.publisher` on the correct topics
 - **Never commit API keys or secrets.**
 - All configs/env variables should be treated as sensitive.
-

FAQ

- **ImportError?** Run `pip install -r requirements.txt` inside the venv.
 - **GSM permissions error?** Double-check service account and IAM roles.
 - **Pub/Sub publish error?** Check topic names and IAM role.
-

Authors & Support

- **Maintainer:** Jithendra Gunuru
 - **Email:** jithendra.gunuru2016@gmail.com
-